EXECUTIVE MEMO

Collections Performance Review

To: Leadership Team From: Garv Kukreja Date: 26 August 2026

Re: Is recovery really up 11% month-on-month, and where should the ₹10 Cr go?

Period reviewed: 1 January – 31 July 2026 (212 days)

BOTTOM LINE

Recovery is essentially flat. The reported 11% month-on-month improvement is almost entirely explained by February having 28 days versus March's 31. I would not deploy the ₹10 Cr yet. Instead, spend ₹1.2 Cr on a 90-day randomized holdout that can establish which collection lever actually works.

1. What happened

The reported 11% improvement is a calendar effect, not evidence of better recovery. February had 28 collection days and March had 31. That difference alone creates a 10.71% increase in monthly collections before any operational change. Once collections are normalized to a daily rate, the underlying improvement is only 0.43%.

Measure	Result
Reported Feb → Mar improvement	+11.19%
Effect of 28 → 31 collection days	+10.71%
Genuine change in daily collection rate	+0.43%
Share explained by calendar	~96%

The decomposition reconciles: 1.1071 × 1.0043 = 1.1119. More broadly, the seven-month daily recovery rate is flat; January to July is down about 1.1%, and the fitted trend is not distinguishable from zero. Across six reasonable definitions of “recovery,” the movement ranges from −1.6% to +0.9%.

- The extract is not a full year. It runs through 8 August 2026, with only eight August days. Including raw August in a monthly dashboard would create an artificial ~74% collapse—the same calendar problem in reverse.
- Reversals matter. ₹9.16 Cr of collections were reversed during the period, equal to 7.23% of gross collections. A gross-only view therefore overstates performance materially.

2. What we tested

The data does not point to an operational explanation for the reported movement.

- Portfolio mix is stable. High-risk share stayed within 24.69%–25.68%, while mean DPD stayed around 55.8–56.8 days. No new accounts were originated during the review window.
- Reach is stable. Accounts attempted ranged from 9,531–10,417 per month and accounts reached from 2,188–2,491. Contact rate remained between 22.50% and 24.02%, with no downward trend.
- Disposition coding is stable. All three schema versions run concurrently across the seven months in broadly similar volumes, so a coding migration does not explain the movement.

The simplest explanation is therefore a measurement artifact: unequal-length months were compared directly, and a single favorable month-pair was used to represent the trend. The other month-pairs do not support the same story.

3. The finding that matters for the ₹10 Cr decision

No operational lever in this dataset has a reliable relationship with payment. That is the key reason I would avoid making a large allocation based on the current evidence.

Test	Result	Read-through
Never touched on any channel	37 / 30,000 (0.12%)	Too small to serve as a useful control group
Never touched → paid	32.43% (95% CI 17.3%–47.5%)	Unstable; overlaps other groups
Dialled, never answered → paid	42.89% (n=16,786)	—
Answered at least once → paid	43.61% (n=13,177)	Reach effect: +0.72 pp
Reach effect	p = 0.212	Not statistically significant
DPD 0–5 vs 90+ recovery/account/day	₹200.09 vs ₹195.12	Only 2.5% difference
LOW vs NPA recovery/account/day	₹201.35 vs ₹197.14	Only 2.1% difference
PTP marked KEPT → payment within 30 days	7.75%	Field is not credible
PTP marked BROKEN → payment within 30 days	6.63%	Barely different from KEPT

Three implications stand out. First, reaching a borrower does not appear to change payment behavior: the observed +0.72 percentage-point difference is noise at this sample size. Second, DPD and risk grade—the two variables we would normally expect to drive collections—show almost no separation. Third, promise-to-pay data is not reliable enough to use for forecasting or targeting.

4. What we can and cannot say with confidence

High confidence: the calendar decomposition is arithmetic and reconciles to two decimal places. The flat-trend result also survives six definitions of recovery. The null findings are reported against effect sizes the tests could detect, rather than simply labeling everything “no effect.”

Important limits: this extract cannot reliably measure the impact of an individual channel, agent, vendor, or calling time. Call volume is uniform across all 24 hours, so time-of-day effects cannot be estimated. Agent identity is not resolvable, so agent/team/vendor comparisons are not supportable. Right-party contact, PTP kept rate, and cost per rupee recovered cannot be computed reliably from the available fields.

Attribution is also highly sensitive to the chosen window: last-touch credit covers only 3.3% of successful payments at a one-day window but 83.4% at 90 days. That 26× swing makes attribution unsuitable for a budget decision without a defined measurement rule.

The premise that targeting changed materially mid-year is not supported. The strongest candidate break date, 25 March, gives $p = 0.061$ before correcting for the 152 candidate dates and $p = 1.00$ after correction. A difference-in-differences test at that date finds ₹+21 per account-week, with a 95% CI of [−₹94, +₹136]. Parallel trends are satisfied ($p = 0.70$), and four placebo dates produce no false positives.

The counterfactual is effectively unchanged: without the targeting change, estimated recovery is ₹116.9 Cr versus ₹117.5 Cr actual, with a confidence interval that includes the actual result.

5. Recommendation

Do not deploy the ₹10 Cr yet. Spend ₹1.2 Cr to learn which lever works, then release the balance against evidence.

Action	Cost	Why
90-day randomized holdout across 60,000 accounts	₹1.2 Cr	Only design here that can answer the causal question the observational data cannot
Repair the PTP status field	~₹15 L	Low-cost fix that removes a major source of bad forecasting
Hold remaining budget pending results	₹8.8 Cr	A one-quarter delay is cheaper than committing ₹8.8 Cr on weak evidence

If leadership requires one area to fund today, I would choose borrower targeting—but explicitly as an economic hypothesis, not a data-backed conclusion. It is the only lever whose cost scales with accounts rather than headcount and is therefore relatively reversible if the test fails.

6. What the economics look like

A 2 percentage-point lift in payment rate would mean roughly 600 additional paying accounts. At an average recovery of ₹90,823 per paying account, that is approximately ₹0.78 Cr per month, or ₹9.3 Cr annually. On the full ₹10 Cr deployment, simple break-even is about 13 months.

The downside of the holdout is that it could be underpowered or contaminated and produce no usable learning. The proposed 60,000-account design is sized to detect a 2 percentage-point difference at 80% power, which is below the effect size the current observational data could not rule out.

7. Reporting changes to make immediately

1. Stop comparing monthly totals. Report collections per collection day so calendar length cannot create a false trend.

2. Report net of reversals. Exclude FAILED and PENDING transactions. Naively summing the amount field overstates recovery by about 35%.

3. Pause PTP kept-rate reporting. Do not use the field for forecasting or performance management until its definition and status transitions are repaired.

8. Bottom line for leadership

The evidence does not show an 11% improvement in recovery. It shows a roughly flat collection rate, a reporting methodology that magnifies calendar differences, and no currently measurable operational lever that deserves ₹10 Cr of incremental spend. The disciplined move is to spend ₹1.2 Cr on a controlled experiment, repair the data needed for measurement, and release the remaining ₹8.8 Cr only once the experiment identifies a causal lift.

Evidence base: 17 forensic tests (13 confirmed, 4 explanations ruled out); a golden dataset whose Raw → Rejected → Golden lineage reconciles exactly for every entity; and reproducible SQL for the quantitative results. Every metric was independently re-derived from the raw CSVs; one metric failed the audit and was corrected before this memo.